




Two decisions worth defending:

`raw_hash` **is verifiable without this system**. An investigator holding the original log file can recompute SHA-256 over a line and find the matching stored event — no ledger, no database, no need to trust us. That is what makes it evidence rather than a checksum.

Clock skew is flagged, not fixed. 141 events in the development corpus are stamped 2012 by an appliance with a dead RTC. Rewriting the timestamp would make the normalised record disagree with the raw line, breaking the losslessness guarantee. Analytics can exclude them; a SOC should arguably be told, because a box with the wrong clock produces evidence that will not correlate.

3. Tamper-evident ledger

Not a blockchain, and saying so would be overclaiming. The primitive underneath one, where it earns its place.

```
event -> leaf  = SHA256(0x00 || canonical(event))
leaves -> root = Merkle tree, odd nodes
                    promoted, never duplicated
block  -> hash  = SHA256(header || merkle_root)
blocks -> chain = each block commits to the
                    hash of the one before it
chain  -> head  = the single value that
                    must be protected
```

Attack	Result
Edit a message	Caught — Merkle mismatch
Reorder or delete events	Caught
Re-seal the block to hide the edit	Caught — breaks the link to the next block
Re-seal the entire chain	Not caught internally — self-consistent; caught only by an externally held head
Forge an inclusion proof	Rejected

- **RFC 6962 domain separation** — without distinct leaf/node prefixes an internal node can be replayed as a leaf (Merkle second-preimage attack).
- **Odd nodes promoted, not duplicated** — duplicating is what made Bitcoin's CVE-2012-2459 possible: two different trees yielding one root.
- **Nine fields sealed**, deliberately excluding `processed_at` and `embedding`, which change legitimately on re-parse. Sealing those would flag honest re-ingestion as tampering, and a control that cries wolf gets switched off.
- **9 sibling hashes** prove one event was in a 500-event block *without revealing the other 499* — which matters when the rest of the log is classified.

Tamper-evident, not tamper-proof. Publishing the head hash somewhere the operator does not control is what turns evidence into proof, and is the one place a real blockchain or an RFC 3161 timestamp authority would belong.

4. Onboarding an unseen source

```
unknown lines
-> tokenize into typed slots
-> group by shape, then align
-> literals = structure
    varying slots = fields
v
parser_semantics: what does each field MEAN?
  position      RFC 3164/5424 fix the host
                  and program slots
  key name       vocabulary from the formats
                  we already parse
  value shape    ipv4, timestamp, severity
v
a runnable detector + per-field evidence
```

Measured on a vendor absent from the codebase: **3 ms**, and the generated detector matched **4/4 lines with zero human edits**. It records the evidence that named each field as a comment, so a reviewer can disagree with one inference rather than distrust the file.

It will not guess whether an unnamed IP is source or destination. Putting wrong evidence into an investigation is the failure this project exists to prevent.

5. Integration and deployment

Pull — streaming exports, flat memory: NDJSON · ECS · OCSF 1.3.0 · CEF · CSV **Push** — syslog (RFC 5424 + RFC 6587 framing, CEF payload) · HTTP NDJSON (Splunk HEC, Elastic, Sentinel) · rotating files for a lake. Bounded queue, exponential backoff, dead-letter counter. Delivery is **at-least-once**; deduplicate on the stable event `id`.

Air-gapped. Six Docker images, ~11.5 GB, `docker load` then `docker compose up`. No runtime downloads, telemetry or licence checks — verified with `docker run --network none`. Speech, vision and document generation all run locally; there is no cloud API anywhere in the path.

Configuration. Each service reads its own `.env`, loaded by `config.py` so every entry point honours it — not just the FastAPI app. Site-specific values live in `.env.example`, never as defaults in code: a framework claiming to be vendor-agnostic should not ship one customer's domain as its built-in assumption.

6. Performance, and its limits

Intel i5-12450H, 8 cores, no GPU.

		Condition
3,513 events/sec	normalisation	1 core, parse stage only
16,826 events/sec	normalisation	8 cores — 5.83× measured
229 events/sec	parse + graph write	semantic search off
3 events/sec	full pipeline	embedding is 94.5% of wall clock

Parsing scales with **CPU cores**, not a GPU — the work is Python's `re` engine, branchy backtracking over short strings, which does not vectorise.

Where this breaks. Neo4j is a correlation and investigation store, not a log archive; at billions of events per day it will not hold, and the streaming exports to a SIEM or data lake are the scale path rather than a nice-to-have. Embedding on CPU is the pipeline bottleneck and is optional per deployment. The live Kafka path has never run against production infrastructure.

We would rather state the boundary than be found at it.